

Programming Assignment #2 -- Multi-Process Matrix Multiplication using Shared Memory

Introduction to Operating Systems

Prof. Li-Pin Chang

CS@NYCU

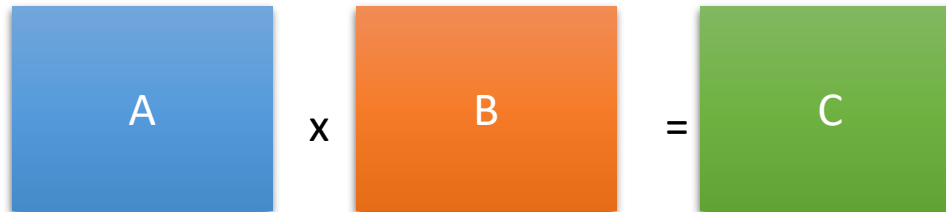
Overview

- Matrix multiplication using multiple processes
 - Basic parallel processing
 - Will be faster on multicore machines
- Input: the dimension of two square matrices A & B
 - E.g., 100 \rightarrow A, B, and C are 100x100 square matrices
- Output: an execution time and a checksum

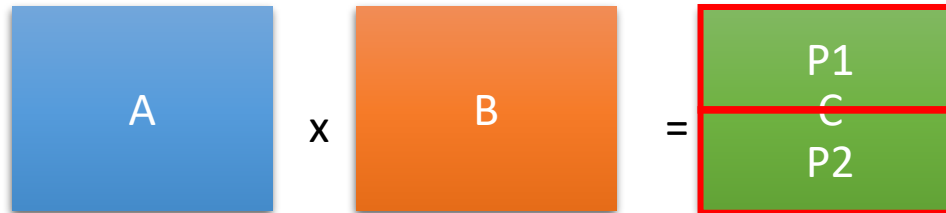


Task Partition

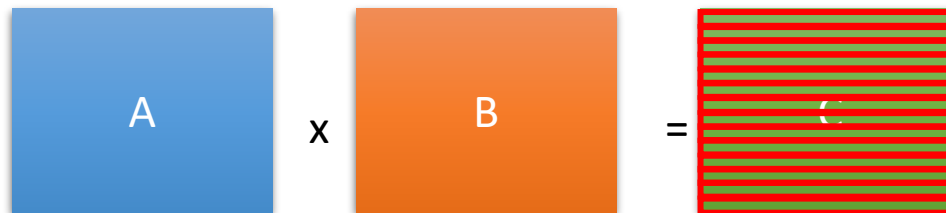
- 1-process matrix multiplication



- 2-process matrix multiplication

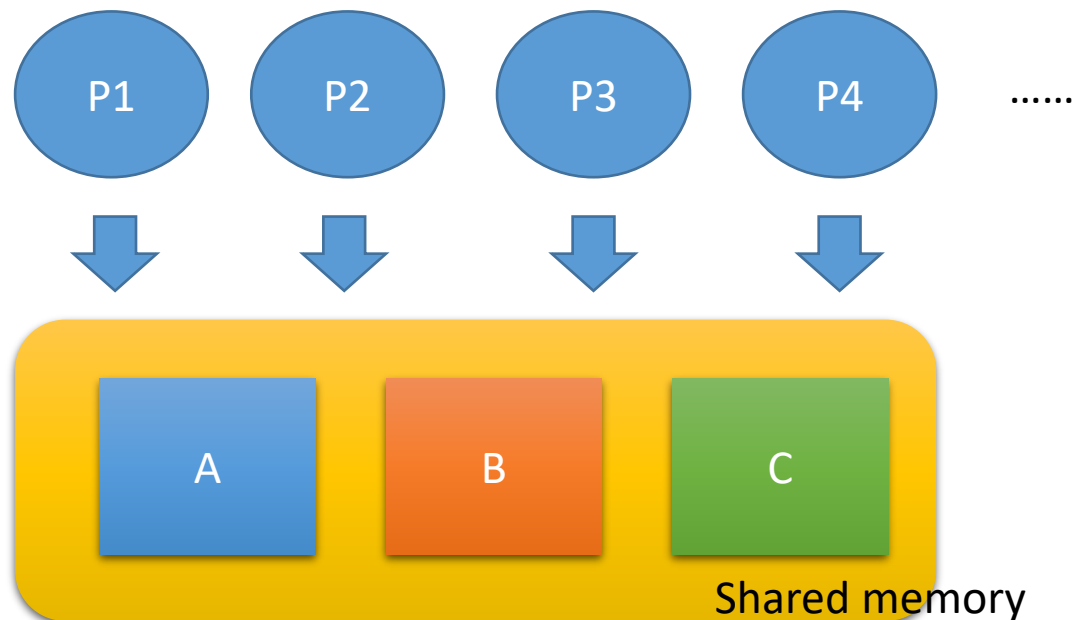


- 16-process matrix multiplication



Shared Memory

- Matrices A, B, and C are stored in a shared memory
- **No** locking/synchronization is required since multiplication on sub-matrices are mutually independent



Header Files

- `unistd.h`
- `sys/ipc.h`
- `sys/shm.h`
- `sys/wait.h`

- `sys/time.h`

APIs

- `shmget()` – create a block of shared memory
- `shmat()` – attach shared memory to the current process's address space
- `shmdt()` – detach shared memory from the current process's address space
- `shmctl()` – control shared memory
- `gettimeofday()`

shmget()

- `int shmget(key_t key, size_t size, int shmflg);`
- create a block of shared memory
- Return the **id of the request shm** of size equals to the value of `size`
- `key` : 0/IPC_PRIVATE for new allocate shm
- `size` : size in bytes
- `shmflg` : IPC_CREAT | mode_flags(9 bits)
e.g. IPC_CREAT | 0600 for read-write shm

shmat()

- `void *shmat(int shmid, const void *shmaddr, int shmflg) ;`
- **Attach** shared memory to the current process's address space
- Return the address of the attached shared memory segment identified by *shmid*
- *shmaddr* : NULL for system to choose suitable address
- *shmflg* : 0 for read/write, **SHM_RDONLY** for read only

shmdt()

- `int shmdt(const void *shmaddr);`
- Detach the shared memory segment located at the address *shmaddr* from the address space of the calling process
- *shmaddr* must equal to the value returned by `shmat()`

shmctl()

- `int shmctl(int shmid, int cmd, struct shmids *buf) ;`
- control shared memory
- Perform control operation
 - `IPC_STAT`
 - `IPC_SET`
 - `IPC_RMID`
 - ...
- `IPC_RMID`: Marking a shared memory to be deleted. The share memory will be destroyed on when the last process detach the memory from its address space. Must be called by the **owner or** creator of the shared memory.
- `[parent]shmget() → shmat()*N → shmdt()*N → [parent]shmctl(RMID)`

gettimeofday()

```
struct timeval start, end;  
gettimeofday(&start, 0);  
//do something  
gettimeofday(&end, 0);  
int sec = end.tv_sec - start.tv_sec;  
int usec = end.tv_usec - start.tv_usec;  
  
printf("elapsed %f sec", sec+(usec/1000000.0));
```

Matrix Initial Values

- Matrix elements in A and B are initialized as follows (for example, an 8x8 matrix)

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39
40	41	42	43	44	45	46	47
48	49	50	51	52	53	54	55
56	57	58	59	60	61	62	63

Checksum of this matrix: 2016

Requirements

- Create worker processes using `fork()`
 - Wait until all worker processes have `exit()`'ed
- About matrices A, B, and C
 - A and B can be inside or outside of the shared memory
 - A and B can be the same or separate matrices
 - C **must** be in the shared memory
- Print the elapsed time and the matrix checksum
 - 16 cases, degree of process parallelism increases from 1 to 16
 - The final checksum must be correct
 - Matrix elements and the checksum are all **32-bit unsigned integers**
 - Don't worry, just let the checksum integer overflow

Requirements – cont'd

- TAs will test your program using any matrix dimension between 100×100 and 800×800
- **[important]** TAs will test your program on a multicore machine
 - Suppose that the test platform has 4 cores, then
 - your 2-process and 3-process versions must be noticeably faster than your 1-process version
 - But the speedup quickly saturates as process # increases
 - **If not, there **must** be something wrong with your program**
- Violating any of these requirements (and those in the prior page) will incur a score penalty

Output

Input the matrix dimension: 800↵

Multiplying matrices using 1 process

Elapsed time: 5.814723 sec, Checksum: 561324032

Multiplying matrices using 2 processes

Elapsed time: 3.10231 sec, Checksum: 561324032

Multiplying matrices using 3 processes

Elapsed time: 2.927338 sec, Checksum: 561324032

.....

Multiplying matrices using 16 processes

Elapsed time: x.xxxxx sec, Checksum: 561324032

API Reference

- <http://blog.csdn.net/guoping16/article/details/6584058>
- <http://man7.org/linux/man-pages/man2/shmget.2.html>
- <http://man7.org/linux/man-pages/man2/shmat.2.html>

Testing OS Environment

- Ubuntu 18.04
- Install as a VM or on a physical machine

Header of your .c or .cpp

```
/*
```

```
Student No.: 31415926
```

```
Student Name: John Doe
```

```
Email: xxx@yyy.zzz
```

```
SE tag: xnxctxuxoxsx
```

```
Statement: I am fully aware that this program is not  
supposed to be posted to a public server, such as a  
public GitHub repository or a public web page.
```

```
*/
```